

TJ Wiegman  
ASM 591 AI  
Lab 7  
2024-11-06

# Examples

## Example 1: Dense Layer and Forward Propagation

In [233...

```
import numpy as np

# Define a simple dense layer
class DenseLayer:
    def __init__(self, input_size, output_size):
        # Randomly initialize weights and biases
        self.weights = np.random.randn(input_size, output_size) * 0.1
        self.biases = np.zeros((1, output_size))

    def forward(self, inputs):
        # Perform forward propagation: inputs * weights + biases
        return np.dot(inputs, self.weights) + self.biases

# Example: 3 inputs, 2 outputs
dense_layer = DenseLayer(3, 2)

# Example input
inputs = np.array([[1.0, 2.0, 3.0]])

# Perform forward propagation
output = dense_layer.forward(inputs)
print("Output of dense layer:", output)
```

Output of dense layer: [[0.29070947 0.16140856]]

## Example 2: Activation Function

In [234...

```
# Define activation functions: ReLU and Sigmoid
def relu(x):
    return np.maximum(0, x)

def sigmoid(x):
    return 1 / (1 + np.exp(-x))

# Example input
inputs = np.array([[1.0, -1.0, 0.0]])

# Applying activation functions
print("ReLU Output:", relu(inputs))
print("Sigmoid Output:", sigmoid(inputs))
```

```
ReLU Output: [[1. 0. 0.]]
Sigmoid Output: [[0.73105858 0.26894142 0.5      ]]
```

## Example 3: Backpropagation

```
In [235... class DenseLayerWithBackprop:
    def __init__(self, input_size, output_size):
        self.weights = np.random.randn(input_size, output_size) * 0.1
        self.biases = np.zeros((1, output_size))

    def forward(self, inputs):
        self.inputs = inputs
        return np.dot(inputs, self.weights) + self.biases

    def backward(self, dvalues, learning_rate=0.001):
        # Gradient on weights and biases
        self.dweights = np.dot(self.inputs.T, dvalues)
        self.dbiases = np.sum(dvalues, axis=0, keepdims=True)

        # Gradient on inputs for chain rule backpropagation
        self.dinputs = np.dot(dvalues, self.weights.T)

        # Update weights and biases using gradient descent
        self.weights -= learning_rate * self.dweights
        self.biases -= learning_rate * self.dbiases

# Forward and Backward Propagation Example
layer = DenseLayerWithBackprop(3, 2)
inputs = np.array([[1.0, 2.0, 3.0]])
output = layer.forward(inputs)

# Example gradient (derivative of some loss function)
dvalues = np.array([[1.0, 1.0]])
layer.backward(dvalues)
```

## Example 4: Feature Handling and Normalization

```
In [236... from sklearn.preprocessing import StandardScaler

# Data example: Each row is a feature vector for a sample
features = np.array([[1, 2], [2, 3], [3, 4], [4, 5]])

# Standardizing features
scaler = StandardScaler()
normalized_features = scaler.fit_transform(features)

print("Normalized Features:", normalized_features)

Normalized Features: [[-1.34164079 -1.34164079]
 [-0.4472136  -0.4472136 ]
 [ 0.4472136   0.4472136 ]
 [ 1.34164079  1.34164079]]
```

## Example 5: Loss Functions

```
In [237... # Mean Squared Error (MSE) Loss function
def mse_loss(y_true, y_pred):
    return np.mean((y_true - y_pred) ** 2)

# Example: Target and prediction
y_true = np.array([[1.0, 0.0]])
y_pred = np.array([[0.8, 0.2]])

# Compute loss
loss = mse_loss(y_true, y_pred)
print("MSE Loss:", loss)
```

MSE Loss: 0.039999999999999994

## Problem Sets

### Problem 1: Build and Train a Neural Network with Forward Pass

Goal:

- Implement a neural network with two dense layers.
- Apply the ReLU activation function after the first layer.
- Apply the Sigmoid activation function after the second layer.
- Write a function that performs the forward pass.

Instructions:

- Implement two layers (using the DenseLayer class provided above).
- Apply the ReLU activation function to the output of the first layer.
- Apply the Sigmoid activation function to the final output.
- Write a forward\_pass function that takes an input, passes it through the two layers, and prints the final output.

```
In [238... class Problem1Net:
    def __init__(self) -> None:
        self.nn1 = DenseLayer(10, 5)
        self.nn2 = DenseLayer(5, 1)

    def forward(self, input):
        x = self.nn1.forward(input)
        x = relu(x)
        x = self.nn2.forward(x)
        return sigmoid(x)

plnet = Problem1Net()
```

```
def forward_pass(input):  
    print(plnet.forward(input))
```

## Problem 2: Implement Backpropagation in a Neural Network

Goal:

- Extend your network to include backpropagation.
- Use the Mean Squared Error (MSE) loss function to calculate the error.
- Implement weight updates using gradient descent.

Instructions:

- Add a backward\_pass method that computes the gradients for the weights and biases.
- After the forward pass, compute the error using the MSE loss function.
- Use backpropagation to update the weights and biases.
- Run multiple iterations and print how the loss decreases over time.

```
In [239... def mse_derivative(y_pred, y_true):  
    return 2*(y_pred - y_true)  
  
def sigmoid_derivative(value):  
    return value * (1 - value)  
  
def relu_derivative(value):  
    return np.where(value > 0, 1, 0)  
  
class Problem2Net:  
    def __init__(self) -> None:  
        self.nn1 = DenseLayerWithBackprop(10, 5)  
        self.nn2 = DenseLayerWithBackprop(5, 1)  
  
    def forward(self, input):  
        self.input = input  
        z1 = self.nn1.forward(input)  
        a1 = relu(z1)  
        z2 = self.nn2.forward(a1)  
        a2 = sigmoid(z2)  
        return a2  
  
    def backward_pass(self, y_true):  
        # Forward pass to get output of each layer  
        z1 = self.nn1.forward(self.input)  
        z2 = self.nn2.forward(relu(z1))  
        a2 = sigmoid(z2)  
  
        # Backpropagate to get gradient deltas  
        d2 = mse_derivative(a2, y_true) * sigmoid_derivative(a2)
```

```

        d1 = np.dot(d2, self.nn2.weights.T) * relu_derivative(z1)

        # Adjust weights and biases with deltas
        self.nn2.backward(d2)
        self.nn1.backward(d1)

p2net = Problem2Net()

# Random demo
x = np.random.random(size=(100,10))
y = np.random.random(size=(100,))

for epoch in range(20):
    losses = []
    for batch in range(x.shape[0]):
        pred = p2net.forward(x[batch, None])
        losses.append(mse_loss(y[batch], pred))
        p2net.backward_pass(y[batch])

    print(f"[{epoch+1}/20] loss = {np.average(loss):.4}")

```

```

[1/20] loss = 0.04
[2/20] loss = 0.04
[3/20] loss = 0.04
[4/20] loss = 0.04
[5/20] loss = 0.04
[6/20] loss = 0.04
[7/20] loss = 0.04
[8/20] loss = 0.04
[9/20] loss = 0.04
[10/20] loss = 0.04
[11/20] loss = 0.04
[12/20] loss = 0.04
[13/20] loss = 0.04
[14/20] loss = 0.04
[15/20] loss = 0.04
[16/20] loss = 0.04
[17/20] loss = 0.04
[18/20] loss = 0.04
[19/20] loss = 0.04
[20/20] loss = 0.04

```

## Problem 3: Feature Handling and Neural Network Implementation

Goal:

- Take a small dataset (like Iris from sklearn).
- Normalize the input features.
- Implement a neural network for classification.

Instructions:

- Use a dataset (you can load the Iris dataset from sklearn.datasets).
- Normalize the features using StandardScaler.
- Build a neural network with two hidden layers using ReLU activation functions.
- Train the network to classify the data points and evaluate its accuracy.

```
In [240... #importing the dataset
from sklearn.datasets import load_iris
data = load_iris()
X = data.data    # Input features
y = data.target # Target labels
print(X.shape, y.shape)
```

(150, 4) (150,)

```
In [241... for i in range(len(y)):
            if y[i] not in y[:i]: print(y[i])
```

0  
1  
2

```
In [242... #Normalize the input features
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_normalized = scaler.fit_transform(X)
```

```
In [243... #splitting dataset into test and train sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_normalized, y, test_size=0.2)
```

```
In [244... # Enable GPU acceleration
import torch
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(device)
```

cuda

```
In [245... # Create NN
import torch.nn as nn
import torch.nn.functional as F

class Problem3Net(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.nn1 = nn.Linear(4, 3) # Linear = fully connected net
        self.nn2 = nn.Linear(3, 2)
        self.nn3 = nn.Linear(2, 1)

    def forward(self, x):
        x = F.relu(self.nn1(x))
        x = F.relu(self.nn2(x))
        x = F.relu(self.nn3(x))
        return x
```

```

In [246... # Create training function
def train(epoch, model, device, optimizer, data, loss_function):
    # Prepare model
    model.to(device)
    model.train()

    for batch_idx, (X, y) in enumerate(data):
        # Load data into GPU
        X = X.to(device)
        y = y.to(device)
        optimizer.zero_grad()

        # Calculate and record output & loss
        output = model(X)
        loss = loss_function(output, y)
        loss.backward()
        optimizer.step()

        # Periodically report on training progress
        print(f"\rEpoch {epoch}: Training {batch_idx * len(X)} / {len(data.dataset)}
              f"(Loss: {loss.item():02.4})", end="")
    print(f"\rEpoch {epoch}: Trained {len(data.dataset)} / {len(data.dataset)} " +
          f"(Loss: {loss.item():02.4})")

```

```

In [247... # Create testing function
def test(epoch, model, device, data, loss_function):
    # Prepare model and data
    model.to(device)
    model.eval()
    test_loss = []
    correct = []

    with torch.no_grad():
        for batch_idx, (X, y) in enumerate(data):
            # Load data into GPU
            X = X.to(device)
            y = y.to(device)

            # Calculate and record output & loss
            output = model(X)
            test_loss.append(loss_function(output, y).item())

            prediction = output.round()
            correct.append(prediction.eq(y).sum() / len(y))

        # Periodically report on testing progress
        print(f"\rEpoch {epoch}: Testing {batch_idx * len(X)} / {len(data.dataset)}
              f"(Loss: {test_loss[-1]:02.4})", end="")
    print(f"\rEpoch {epoch}: Testing {len(data.dataset)} / {len(data.dataset)}")

    # Report results
    test_loss = torch.mean(torch.tensor(test_loss))
    accuracy = 100 * torch.mean(torch.tensor(correct))
    print(f"Test Result, epoch {epoch}: Avg loss {test_loss:04.4}, Accuracy {accuracy:04.4}")

    return accuracy

```

```

In [248... # Create data loaders
from torch.utils.data import Dataset, DataLoader

class IrisSet(Dataset):
    def __init__(self, X, y):
        self.X = X
        self.y = y

    def __len__(self):
        return len(self.y)

    def __getitem__(self, idx):
        return torch.tensor(self.X[idx, :], dtype=torch.float), torch.tensor([self.y[idx]])

train_set = IrisSet(X_train, y_train)
test_set = IrisSet(X_test, y_test)

train_data = DataLoader(
    dataset = train_set,
    batch_size = 10,
    shuffle = True
)

test_data = DataLoader(
    dataset = test_set,
    batch_size = 15,
    shuffle = True
)

```

```

In [258... # Train the network
p3net = Problem3Net()
optimizer = torch.optim.Adam(p3net.parameters(), lr=0.01)
accuracy = 0
i = 0

while accuracy < 95:
    i += 1
    train(
        epoch=i,
        model=p3net,
        device=device,
        optimizer=optimizer,
        data=train_data,
        loss_function=F.mse_loss
    )
    accuracy = test(
        epoch=i,
        model=p3net,
        device=device,
        data=test_data,
        loss_function=F.mse_loss
    )
    if i >= 100: break

```



Epoch 1: Trained 120/120 (Loss: 0.5396))  
Epoch 1: Testing 30/30...  
Test Result, epoch 1: Avg loss 0.3875, Accuracy 53.33%  
Epoch 2: Trained 120/120 (Loss: 0.3937))  
Epoch 2: Testing 30/30...  
Test Result, epoch 2: Avg loss 0.2667, Accuracy 66.67%  
Epoch 3: Trained 120/120 (Loss: 0.2646))  
Epoch 3: Testing 30/30...  
Test Result, epoch 3: Avg loss 0.1923, Accuracy 66.67%  
Epoch 4: Trained 120/120 (Loss: 0.1913))  
Epoch 4: Testing 30/30...  
Test Result, epoch 4: Avg loss 0.1488, Accuracy 66.67%  
Epoch 5: Trained 120/120 (Loss: 0.1164))  
Epoch 5: Testing 30/30...  
Test Result, epoch 5: Avg loss 0.1304, Accuracy 90.0%  
Epoch 6: Trained 120/120 (Loss: 0.08258))  
Epoch 6: Testing 30/30...  
Test Result, epoch 6: Avg loss 0.1209, Accuracy 86.67%  
Epoch 7: Trained 120/120 (Loss: 0.1157))  
Epoch 7: Testing 30/30...  
Test Result, epoch 7: Avg loss 0.09112, Accuracy 93.33%  
Epoch 8: Trained 120/120 (Loss: 0.11))97)  
Epoch 8: Testing 30/30...  
Test Result, epoch 8: Avg loss 0.1047, Accuracy 86.67%  
Epoch 9: Trained 120/120 (Loss: 0.04695))  
Epoch 9: Testing 30/30...  
Test Result, epoch 9: Avg loss 0.07993, Accuracy 90.0%  
Epoch 10: Trained 120/120 (Loss: 0.03784))  
Epoch 10: Testing 30/30...  
Test Result, epoch 10: Avg loss 0.0779, Accuracy 90.0%  
Epoch 11: Trained 120/120 (Loss: 0.07931))  
Epoch 11: Testing 30/30...  
Test Result, epoch 11: Avg loss 0.07379, Accuracy 93.33%  
Epoch 12: Trained 120/120 (Loss: 0.04865))  
Epoch 12: Testing 30/30...  
Test Result, epoch 12: Avg loss 0.06854, Accuracy 93.33%  
Epoch 13: Trained 120/120 (Loss: 0.0378))  
Epoch 13: Testing 30/30...  
Test Result, epoch 13: Avg loss 0.06522, Accuracy 93.33%  
Epoch 14: Trained 120/120 (Loss: 0.0127))  
Epoch 14: Testing 30/30...  
Test Result, epoch 14: Avg loss 0.06647, Accuracy 93.33%  
Epoch 15: Trained 120/120 (Loss: 0.01211))  
Epoch 15: Testing 30/30...  
Test Result, epoch 15: Avg loss 0.06411, Accuracy 93.33%  
Epoch 16: Trained 120/120 (Loss: 0.01828))  
Epoch 16: Testing 30/30...  
Test Result, epoch 16: Avg loss 0.08759, Accuracy 90.0%  
Epoch 17: Trained 120/120 (Loss: 0.08471))  
Epoch 17: Testing 30/30...  
Test Result, epoch 17: Avg loss 0.04913, Accuracy 93.33%  
Epoch 18: Trained 120/120 (Loss: 0.03262))  
Epoch 18: Testing 30/30...  
Test Result, epoch 18: Avg loss 0.09003, Accuracy 90.0%  
Epoch 19: Trained 120/120 (Loss: 0.04557))  
Epoch 19: Testing 30/30...

Test Result, epoch 19: Avg loss 0.05855, Accuracy 93.33%  
Epoch 20: Trained 120/120 (Loss: 0.05961))  
Epoch 20: Testing 30/30...  
Test Result, epoch 20: Avg loss 0.06011, Accuracy 90.0%  
Epoch 21: Trained 120/120 (Loss: 0.05596))  
Epoch 21: Testing 30/30...  
Test Result, epoch 21: Avg loss 0.05929, Accuracy 90.0%  
Epoch 22: Trained 120/120 (Loss: 0.03487))  
Epoch 22: Testing 30/30...  
Test Result, epoch 22: Avg loss 0.0561, Accuracy 93.33%  
Epoch 23: Trained 120/120 (Loss: 0.03216))  
Epoch 23: Testing 30/30...  
Test Result, epoch 23: Avg loss 0.06535, Accuracy 90.0%  
Epoch 24: Trained 120/120 (Loss: 0.07382))  
Epoch 24: Testing 30/30...  
Test Result, epoch 24: Avg loss 0.05702, Accuracy 90.0%  
Epoch 25: Trained 120/120 (Loss: 0.0634))  
Epoch 25: Testing 30/30...  
Test Result, epoch 25: Avg loss 0.04981, Accuracy 93.33%  
Epoch 26: Trained 120/120 (Loss: 0.0475))  
Epoch 26: Testing 30/30...  
Test Result, epoch 26: Avg loss 0.1016, Accuracy 90.0%  
Epoch 27: Trained 120/120 (Loss: 0.06247))  
Epoch 27: Testing 30/30...  
Test Result, epoch 27: Avg loss 0.04751, Accuracy 93.33%  
Epoch 28: Trained 120/120 (Loss: 0.01993))  
Epoch 28: Testing 30/30...  
Test Result, epoch 28: Avg loss 0.07638, Accuracy 90.0%  
Epoch 29: Trained 120/120 (Loss: 0.03991))  
Epoch 29: Testing 30/30...  
Test Result, epoch 29: Avg loss 0.05103, Accuracy 93.33%  
Epoch 30: Trained 120/120 (Loss: 0.04166))  
Epoch 30: Testing 30/30...  
Test Result, epoch 30: Avg loss 0.06542, Accuracy 90.0%  
Epoch 31: Trained 120/120 (Loss: 0.03126))  
Epoch 31: Testing 30/30...  
Test Result, epoch 31: Avg loss 0.05291, Accuracy 90.0%  
Epoch 32: Trained 120/120 (Loss: 0.03266))  
Epoch 32: Testing 30/30...  
Test Result, epoch 32: Avg loss 0.06083, Accuracy 90.0%  
Epoch 33: Trained 120/120 (Loss: 0.006915))  
Epoch 33: Testing 30/30...  
Test Result, epoch 33: Avg loss 0.0631, Accuracy 90.0%  
Epoch 34: Trained 120/120 (Loss: 0.03734))  
Epoch 34: Testing 30/30...  
Test Result, epoch 34: Avg loss 0.04868, Accuracy 93.33%  
Epoch 35: Trained 120/120 (Loss: 0.04058))  
Epoch 35: Testing 30/30...  
Test Result, epoch 35: Avg loss 0.09283, Accuracy 90.0%  
Epoch 36: Trained 120/120 (Loss: 0.03795))  
Epoch 36: Testing 30/30...  
Test Result, epoch 36: Avg loss 0.04474, Accuracy 93.33%  
Epoch 37: Trained 120/120 (Loss: 0.02592))  
Epoch 37: Testing 30/30...  
Test Result, epoch 37: Avg loss 0.1031, Accuracy 90.0%  
Epoch 38: Trained 120/120 (Loss: 0.05173))

```

Epoch 38: Testing 30/30...
Test Result, epoch 38: Avg loss 0.04569, Accuracy 93.33%
Epoch 39: Trained 120/120 (Loss: 0.03547))
Epoch 39: Testing 30/30...
Test Result, epoch 39: Avg loss 0.07273, Accuracy 90.0%
Epoch 40: Trained 120/120 (Loss: 0.04601))
Epoch 40: Testing 30/30...
Test Result, epoch 40: Avg loss 0.06892, Accuracy 90.0%
Epoch 41: Trained 120/120 (Loss: 0.01077))
Epoch 41: Testing 30/30...
Test Result, epoch 41: Avg loss 0.06458, Accuracy 90.0%
Epoch 42: Trained 120/120 (Loss: 0.01762))
Epoch 42: Testing 30/30...
Test Result, epoch 42: Avg loss 0.0491, Accuracy 96.67%

```

## Problem 4: Explore Different Activation and Loss Functions

Goal:

- Experiment with different activation and loss functions.
- Compare the performance of each.

Instructions:

- Modify your previous network to use Leaky ReLU and softmax activation functions.
- Implement two new loss functions:
  - Categorical Cross-Entropy for classification.
  - Hinge Loss for multi-class classification.
- Compare the training accuracy and loss with each activation and loss function.

```

In [250...] import torch.nn.functional as F

class Problem4Net(nn.Module):
    def __init__(self, f_activation) -> None:
        super().__init__()
        self.act = f_activation
        self.nn1 = nn.Linear(4, 3) # Linear = fully connected net
        self.nn2 = nn.Linear(3, 2)
        self.nn3 = nn.Linear(2, 1)

    def forward(self, x):
        x = self.act(self.nn1(x))
        x = self.act(self.nn2(x))
        x = self.act(self.nn3(x))
        return x

leakyModel = Problem4Net(F.leaky_relu)
softmModel = Problem4Net(F.softmax)

```

```

In [251...] MAX_EPOCHS = 5

```

```

for model in [leakyModel, softmModel]:
    for loss_func in [F.cross_entropy, F.hinge_embedding_loss]:
        print("="*50)
        print(f"Testing {'Leaky ReLU' if model == leakyModel else 'SoftMax'} Activ
              f"{'Cross-Entropy' if loss_func == F.cross_entropy else 'Hinge'} Los
        print("-"*50)
        optimizer = torch.optim.Adam(model.parameters())
        for i in range(1, MAX_EPOCHS+1):
            train(
                epoch=i,
                model=model,
                device=device,
                optimizer=optimizer,
                data = train_data,
                loss_function=loss_func
            )
            test(
                epoch=i,
                model=model,
                device=device,
                data=test_data,
                loss_function=loss_func
            )
        print("="*50)

```

=====  
Testing Leaky ReLU Activation + Cross-Entropy Loss Function:  
-----

Epoch 1: Trained 120/120 (Loss: -0.0))  
Epoch 1: Testing 30/30...  
Test Result, epoch 1: Avg loss 00.0, Accuracy 50.0%  
Epoch 2: Trained 120/120 (Loss: -0.0))  
Epoch 2: Testing 30/30...  
Test Result, epoch 2: Avg loss 00.0, Accuracy 50.0%  
Epoch 3: Trained 120/120 (Loss: -0.0))  
Epoch 3: Testing 30/30...  
Test Result, epoch 3: Avg loss 00.0, Accuracy 50.0%  
Epoch 4: Trained 120/120 (Loss: -0.0))  
Epoch 4: Testing 30/30...  
Test Result, epoch 4: Avg loss 00.0, Accuracy 50.0%  
Epoch 5: Trained 120/120 (Loss: -0.0))  
Epoch 5: Testing 30/30...  
Test Result, epoch 5: Avg loss 00.0, Accuracy 50.0%

=====  
Testing Leaky ReLU Activation + Hinge Loss Function:  
-----

Epoch 1: Trained 120/120 (Loss: 0.9174))  
Epoch 1: Testing 30/30...  
Test Result, epoch 1: Avg loss 0.9172, Accuracy 50.0%  
Epoch 2: Trained 120/120 (Loss: 0.9084))  
Epoch 2: Testing 30/30...  
Test Result, epoch 2: Avg loss 0.9028, Accuracy 50.0%  
Epoch 3: Trained 120/120 (Loss: 0.9073))  
Epoch 3: Testing 30/30...  
Test Result, epoch 3: Avg loss 0.8885, Accuracy 50.0%  
Epoch 4: Trained 120/120 (Loss: 0.9173))  
Epoch 4: Testing 30/30...  
Test Result, epoch 4: Avg loss 0.8748, Accuracy 50.0%  
Epoch 5: Trained 120/120 (Loss: 0.9359))  
Epoch 5: Testing 30/30...  
Test Result, epoch 5: Avg loss 0.8618, Accuracy 50.0%

=====  
=====  
Testing SoftMax Activation + Cross-Entropy Loss Function:  
-----

Epoch 1: Trained 120/120 (Loss: -0.0))  
Epoch 1: Testing 30/30...  
Test Result, epoch 1: Avg loss 00.0, Accuracy 50.0%  
Epoch 2: Trained 120/120 (Loss: -0.0))  
Epoch 2: Testing 30/30...  
Test Result, epoch 2: Avg loss 00.0, Accuracy 50.0%  
Epoch 3: Trained 120/120 (Loss: -0.0))  
Epoch 3: Testing 30/30...  
Test Result, epoch 3: Avg loss 00.0, Accuracy 50.0%  
Epoch 4: Trained 120/120 (Loss: -0.0))  
Epoch 4: Testing 30/30...  
Test Result, epoch 4: Avg loss 00.0, Accuracy 50.0%  
Epoch 5: Trained 120/120 (Loss: -0.0))  
Epoch 5: Testing 30/30...  
Test Result, epoch 5: Avg loss 00.0, Accuracy 50.0%

Testing SoftMax Activation + Hinge Loss Function:

```
-----  
Epoch 1: Trained 120/120 (Loss: 1.0))  
Epoch 1: Testing 30/30...  
Test Result, epoch 1: Avg loss 01.0, Accuracy 50.0%  
Epoch 2: Trained 120/120 (Loss: 1.0))  
Epoch 2: Testing 30/30...  
Test Result, epoch 2: Avg loss 01.0, Accuracy 50.0%  
Epoch 3: Trained 120/120 (Loss: 1.0))  
Epoch 3: Testing 30/30...  
Test Result, epoch 3: Avg loss 01.0, Accuracy 50.0%  
Epoch 4: Trained 120/120 (Loss: 1.0))  
Epoch 4: Testing 30/30...  
Test Result, epoch 4: Avg loss 01.0, Accuracy 50.0%  
Epoch 5: Trained 120/120 (Loss: 1.0))  
Epoch 5: Testing 30/30...  
Test Result, epoch 5: Avg loss 01.0, Accuracy 50.0%  
=====
```

```
/tmp/ipykernel_2284320/2914991766.py:12: UserWarning: Implicit dimension choice for  
softmax has been deprecated. Change the call to include dim=X as an argument.  
  x = self.act(self.nn1(x))  
/tmp/ipykernel_2284320/2914991766.py:13: UserWarning: Implicit dimension choice for  
softmax has been deprecated. Change the call to include dim=X as an argument.  
  x = self.act(self.nn2(x))  
/tmp/ipykernel_2284320/2914991766.py:14: UserWarning: Implicit dimension choice for  
softmax has been deprecated. Change the call to include dim=X as an argument.  
  x = self.act(self.nn3(x))
```

## Problem 5: Build a Neural Network for Regression

Goal:

- Implement a neural network to predict continuous values.
- Train the network on a dataset like Boston Housing Prices.

Instructions:

- Use the Boston Housing dataset from sklearn.
- Normalize both the features and target values.
- Implement a network with two hidden layers, using ReLU as the activation.
- Train the network using MSE as the loss function.
- Plot the predicted vs. actual values after training.

```
In [252... from sklearn.datasets import load_boston
```

```

-----
ImportError                                Traceback (most recent call last)
/tmp/ipykernel_2284320/2305181818.py in <module>
----> 1 from sklearn.datasets import load_boston

~/local/lib/python3.10/site-packages/sklearn/datasets/__init__.py in __getattr__(name)
    154         """
    155     )
--> 156     raise ImportError(msg)
    157     try:
    158         return globals()[name]

```

**ImportError:**  
`load\_boston` has been removed from scikit-learn since version 1.2.

The Boston housing prices dataset has an ethical problem: as investigated in [1], the authors of this dataset engineered a non-invertible variable "B" assuming that racial self-segregation had a positive impact on house prices [2]. Furthermore the goal of the research that led to the creation of this dataset was to study the impact of air quality but it did not give adequate demonstration of the validity of this assumption.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```

import pandas as pd
import numpy as np

data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]

```

Alternative datasets include the California housing dataset and the Ames housing dataset. You can load the datasets as follows::

```

from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()

```

for the California housing dataset and::

```

from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)

```

for the Ames housing dataset.

[1] M Carlisle.

"Racist data destruction?"

<<https://medium.com/@docintangible/racist-data-destruction-113e3eff54a8>>

```
[2] Harrison Jr, David, and Daniel L. Rubinfeld.  
"Hedonic housing prices and the demand for clean air."  
Journal of environmental economics and management 5.1 (1978): 81-102.  
<https://www.researchgate.net/publication/4974606\_Hedonic\_housing\_prices\_and\_the\_demand\_for\_clean\_air>
```

```
In [406... from sklearn.datasets import fetch_openml  
housing = fetch_openml(name="house_prices", as_frame=True)  
  
# Preview the inputs  
housing.data.head()
```

```
Out[406...   Id  MSSubClass  MSZoning  LotFrontage  LotArea  Street  Alley  LotShape  LandContour  Utilities  
0    1         60      RL         65.0     8450   Pave   NaN     Reg         Lvl     AllPub  
1    2         20      RL         80.0     9600   Pave   NaN     Reg         Lvl     AllPub  
2    3         60      RL         68.0    11250   Pave   NaN     IR1         Lvl     AllPub  
3    4         70      RL         60.0     9550   Pave   NaN     IR1         Lvl     AllPub  
4    5         60      RL         84.0    14260   Pave   NaN     IR1         Lvl     AllPub
```

5 rows × 80 columns

```
In [407... # Preview the targets  
housing.target.head()
```

```
Out[407... 0    208500  
1    181500  
2    223500  
3    140000  
4    250000  
Name: SalePrice, dtype: int64
```

```
In [408... print(housing.data.shape)  
  
(1460, 80)
```

```
In [409... unique_types = []  
for key in housing.data.keys():  
    t = type(housing.data[key][0])  
    if t not in unique_types: unique_types.append(t)  
  
print(unique_types)  
  
[<class 'numpy.int64'>, <class 'str'>, <class 'numpy.float64'>, <class 'float'>]
```

```
In [410... numeric_types = []  
for key in housing.data:  
    numeric_types.append(key)  
    for val in housing.data[key]:  
        if type(val) not in [np.int64, np.float64, float, int]:  
            if key in numeric_types: numeric_types.remove(key)  
numeric_types.remove("Id")  
print(f"{len(numeric_types)} numeric variables per listing: {numeric_types}")
```



36 numeric variables per listing: ['MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces', 'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal', 'MoSold', 'YrSold']

```
In [411... for key in housing.data.keys():
    if type(housing.data[key][0]) == str:
        unique_vals = []
        for item in housing.data[key]:
            if item not in unique_vals: unique_vals.append(item)

        print(f"{key} can be any of {unique_vals}")
```

MSZoning can be any of ['RL', 'RM', "'C (all)'", 'FV', 'RH']  
 Street can be any of ['Pave', 'Grvl']  
 LotShape can be any of ['Reg', 'IR1', 'IR2', 'IR3']  
 LandContour can be any of ['Lvl', 'Bnk', 'Low', 'HLS']  
 Utilities can be any of ['AllPub', 'NoSeWa']  
 LotConfig can be any of ['Inside', 'FR2', 'Corner', 'CulDSac', 'FR3']  
 LandSlope can be any of ['Gtl', 'Mod', 'Sev']  
 Neighborhood can be any of ['CollgCr', 'Veenker', 'Crawfor', 'NoRidge', 'Mitchel',  
 'Somerst', 'NWAmes', 'OldTown', 'BrkSide', 'Sawyer', 'NridgHt', 'NAMES', 'SawyerW',  
 'IDOTRR', 'MeadowV', 'Edwards', 'Timber', 'Gilbert', 'StoneBr', 'ClearCr', 'NPKvill',  
 'Blmngtn', 'BrDale', 'SWISU', 'Blueste']  
 Condition1 can be any of ['Norm', 'Feedr', 'PosN', 'Artery', 'RRAe', 'RRNn', 'RRAn',  
 'PosA', 'RRNe']  
 Condition2 can be any of ['Norm', 'Artery', 'RRNn', 'Feedr', 'PosN', 'PosA', 'RRAn',  
 'RRAe']  
 BldgType can be any of ['1Fam', '2fmCon', 'Duplex', 'TwnhsE', 'Twnhs']  
 HouseStyle can be any of ['2Story', '1Story', '1.5Fin', '1.5Unf', 'SFoyer', 'SLvl',  
 '2.5Unf', '2.5Fin']  
 RoofStyle can be any of ['Gable', 'Hip', 'Gambrel', 'Mansard', 'Flat', 'Shed']  
 RoofMatl can be any of ['CompShg', 'WdShngl', 'Metal', 'WdShake', 'Membran', 'Tar&G  
 rv', 'Roll', 'ClyTile']  
 Exterior1st can be any of ['VinylSd', 'MetalSd', "'Wd Sdng'", 'HdBoard', 'BrkFace',  
 'WdShng', 'CemntBd', 'Plywood', 'AsbShng', 'Stucco', 'BrkComm', 'AsphShn', 'Stone',  
 'ImStucc', 'CBlock']  
 Exterior2nd can be any of ['VinylSd', 'MetalSd', "'Wd Shng'", 'HdBoard', 'Plywood',  
 "'Wd Sdng'", 'CmentBd', 'BrkFace', 'Stucco', 'AsbShng', "'Brk Cmn'", 'ImStucc', 'As  
 phShn', 'Stone', 'Other', 'CBlock']  
 MasVnrType can be any of ['BrkFace', 'None', 'Stone', 'BrkCmn', nan]  
 ExterQual can be any of ['Gd', 'TA', 'Ex', 'Fa']  
 ExterCond can be any of ['TA', 'Gd', 'Fa', 'Po', 'Ex']  
 Foundation can be any of ['PConc', 'CBlock', 'BrkTil', 'Wood', 'Slab', 'Stone']  
 BsmtQual can be any of ['Gd', 'TA', 'Ex', nan, 'Fa']  
 BsmtCond can be any of ['TA', 'Gd', nan, 'Fa', 'Po']  
 BsmtExposure can be any of ['No', 'Gd', 'Mn', 'Av', nan]  
 BsmtFinType1 can be any of ['GLQ', 'ALQ', 'Unf', 'Rec', 'BLQ', nan, 'LwQ']  
 BsmtFinType2 can be any of ['Unf', 'BLQ', nan, 'ALQ', 'Rec', 'LwQ', 'GLQ']  
 Heating can be any of ['GasA', 'GasW', 'Grav', 'Wall', 'OthW', 'Floor']  
 HeatingQC can be any of ['Ex', 'Gd', 'TA', 'Fa', 'Po']  
 CentralAir can be any of ['Y', 'N']  
 Electrical can be any of ['SBrkr', 'FuseF', 'FuseA', 'FuseP', 'Mix', nan]  
 KitchenQual can be any of ['Gd', 'TA', 'Ex', 'Fa']  
 Functional can be any of ['Typ', 'Min1', 'Maj1', 'Min2', 'Mod', 'Maj2', 'Sev']  
 GarageType can be any of ['Attchd', 'Detchd', 'BuiltIn', 'CarPort', nan, 'Basment',  
 '2Types']  
 GarageFinish can be any of ['RFn', 'Unf', 'Fin', nan]  
 GarageQual can be any of ['TA', 'Fa', 'Gd', nan, 'Ex', 'Po']  
 GarageCond can be any of ['TA', 'Fa', nan, 'Gd', 'Po', 'Ex']  
 PavedDrive can be any of ['Y', 'N', 'P']  
 SaleType can be any of ['WD', 'New', 'COD', 'ConLD', 'ConLI', 'CWD', 'ConLw', 'Con',  
 'Oth']  
 SaleCondition can be any of ['Normal', 'Abnorml', 'Partial', 'AdjLand', 'Alloca',  
 'Family']

In [412... *# Encoding all those str variables as numbers would be a lot of work*  
*# and it's 4am and I've been working on this all night so I'm just going to not*

```
In [578... # Create NN
import torch.nn as nn
import torch.nn.functional as F

class Problem5Net(nn.Module):
    def __init__(self) -> None:
        super().__init__()
        self.nn1 = nn.Linear(len(numeric_types), 16)
        self.nn2 = nn.Linear(16, 16)
        self.nn3 = nn.Linear(16, 1)

    def forward(self, x):
        x = x.to(torch.float32)
        x = F.relu(self.nn1(x))
        x = F.relu(self.nn2(x))
        x = F.relu(self.nn3(x))
        return x
```

```
In [579... class CustomScaler():
    def fit_transform(self, tensor):
        self.offset = torch.mean(tensor)
        self.var = torch.var(tensor)
        return (tensor - self.offset) / self.var

    def inverse_transform(self, tensor):
        return (self.var * tensor) + self.offset
```

```
In [580... # Create data loaders
from torch.utils.data import Dataset, DataLoader
from random import sample

class HouseSet(Dataset):
    def __init__(self, data, targets, ids):
        self.X = torch.zeros(size=(len(ids), len(numeric_types)), dtype=torch.float)
        self.y = np.array(targets[ids]).reshape(1, -1)
        self.scalers = []
        self.vars = []

        # Normalize numeric features
        i = 0
        for key in numeric_types:
            self.scalers.append(CustomScaler())
            self.vars.append(key)
            x = np.nan_to_num(np.array(data[key][ids]), copy=False)
            self.X[:, i] = self.scalers[-1].fit_transform(torch.tensor(x, dtype=torch.float))
            i += 1

        # Normalize targets too
        self.scalers.append(CustomScaler())
        self.y = self.scalers[-1].fit_transform(torch.tensor(self.y, dtype=torch.float))
        self.vars.append("Target")

    def __len__(self):
        return self.X.shape[0]
```

```

    def __getitem__(self, idx):
        return self.X[idx, :], self.y[:, idx]

# Randomly split training/test data
N = housing.data.shape[0]
testN = int(0.2*N)
test_ids = sample(range(N), testN)
train_ids = [i for i in range(N) if i not in test_ids]

# Create datasets
train_set5 = HouseSet(
    data = housing.data,
    targets = housing.target,
    ids = train_ids
)
test_set5 = HouseSet(
    data = housing.data,
    targets = housing.target,
    ids = test_ids
)

# Create data loaders
train_data5 = DataLoader(
    dataset = train_set5,
    batch_size = 5,
    shuffle = True
)
test_data5 = DataLoader(
    dataset = test_set5,
    batch_size = 20,
    shuffle = True
)

```

```

In [581]: # Create testing function
def test(epoch, model, device, data, loss_function):
    # Prepare model and data
    model.to(device)
    model.eval()
    test_loss = []
    predicted = torch.tensor([[0]], device=device)
    real = torch.tensor([[0]], device=device)

    with torch.no_grad():
        for batch_idx, (X, y) in enumerate(data):
            # Load data into GPU
            X = X.to(device)
            y = y.to(device)

            # Calculate and record output & loss
            output = model(X)
            test_loss.append(loss_function(output, y).item())

            predicted = torch.cat((predicted, output))
            real = torch.cat((real, y))

    # Periodically report on testing progress

```

```

        print(f"\rEpoch {epoch}: Testing {batch_idx*len(X)}/{len(data.dataset)}")
        print(f"\rEpoch {epoch}: Testing {len(data.dataset)}/{len(data.dataset)}")

    # Report results
    test_loss = torch.mean(torch.tensor(test_loss))
    print(f"Test Result, epoch {epoch}: Avg loss {test_loss:04.4}")

    real = data.dataset.scalers[-1].inverse_transform(real)
    predicted = data.dataset.scalers[-1].inverse_transform(predicted)
    return real, predicted

```

```

In [582... p5net = Problem5Net()
optimizer = torch.optim.Adam(p5net.parameters())

preds = []
reals = []

for epoch in range(1, MAX_EPOCHS+1):
    train(
        epoch = epoch,
        model = p5net,
        device = device,
        optimizer = optimizer,
        data = train_data5,
        loss_function = F.mse_loss
    )

    real, predicted = test(
        epoch = epoch,
        model = p5net,
        device = device,
        data = test_data5,
        loss_function = F.mse_loss
    )
    preds.append(predicted.cpu())
    reals.append(real.cpu())

```

```

Epoch 1: Trained 1168/1168 (Loss: 1.139e-10)
Epoch 1: Testing 292/292...
Test Result, epoch 1: Avg loss 1.591e-10
Epoch 2: Trained 1168/1168 (Loss: 5.197e-11)
Epoch 2: Testing 292/292...
Test Result, epoch 2: Avg loss 1.56e-10
Epoch 3: Trained 1168/1168 (Loss: 2.44e-10)0)
Epoch 3: Testing 292/292...
Test Result, epoch 3: Avg loss 1.559e-10
Epoch 4: Trained 1168/1168 (Loss: 2.331e-10)
Epoch 4: Testing 292/292...
Test Result, epoch 4: Avg loss 1.561e-10
Epoch 5: Trained 1168/1168 (Loss: 5.988e-11)
Epoch 5: Testing 292/292...
Test Result, epoch 5: Avg loss 1.577e-10

```

```

In [587... import matplotlib.pyplot as plt

plt.figure(figsize=(8,8))

```

```

for i in range(len(reals)):
    plt.plot(reals[i], label=f"True Values #{i+1}")
    plt.plot(preds[i], label=f"Predicted for #{i+1}")
plt.legend()

```

Out[587... <matplotlib.legend.Legend at 0x751fc0077b20>

